

PROBLEM 1.

```
3 #standardizing data X
4 standardizeX = function(X)
5 {
6   mX = matrix(colMeans(X), nr=dim(X)[1], nc=dim(X)[2], byrow=T)
7   X0 = X - mX
8   sd = sqrt(colMeans(X0^2))
9   X0 = X0%%diag(1/sd)
10  output = list(X0=X0, diag=sd)
11  return(output)
12 }
13
14 #subgradient descent for optimizing linear MAE
15 sub_gd = function(X,y,eta)
16 {
17   n = dim(X)[1]
18   d = dim(X)[2]
19   standX = standardizeX(X)
20   X0 = cbind(rep(1,n), standX$X0)
21
22   b0 = numeric(d+1)
23   res0 = y
24   b1 = rep(1, d+1)
25   iter = 0
26
27   repeat{
28     if(max(abs(b1-b0)) < 10^(-5) | iter > 20000) break
29
30     b1 = b0
31     subgrad = sign(res0)
32     grad = -t(X0) %*% subgrad/n
33     b0 = b0 - eta*grad
34     res0 = y - X0%%b0
35     iter = iter + 1
36   }
37
38   b0[-1] = b0[-1]/standX$diag
39   b0[1] = b0[1] - sum(colMeans(X)*b0[-1])
40   output = list(beta=b0, res=res0, iter=iter)
41   return(output)
42 }
43 }
```

PROBLEM 2.

```
48 df = read.csv("education.csv")
49 X = data.matrix(df[,4:6])
50 y = data.matrix(df[,7])
51 #subgradient descent for MAE
52 maefit <- sub_gd(X,y,1)
53 #MSE
54 lsefit <- lm(y~X)
```

MAE estimator output:

```
$beta
      [,1]
[1,] -356.75802116
[2,]  0.02870102
[3,]  0.06467542
[4,]  0.96283273
```

MSE estimator output:

```
Coefficients:
(Intercept)      xx1      xx2      xx3
-5.566e+02  -4.269e-03  7.239e-02  1.552e+00
```

In general, MAE tends to select the smaller values (in absolute value).

For example, β_0 for MAE is $|-356.76| < |-556.6|$ for MSE.

PROBLEM 3.

Part 1:


```
59 build_poly <- function(x,d)
60 {
61   X = matrix(nrow=length(x), ncol=d)
62   for(j in 1:d)
63   {
64     X[,j] = x^(j)
65   }
66   X = cbind(1, X)
67   return(X)
68 }
```

Part 2:

```
71 df = read.csv("polynomial1.csv")
72 x = data.matrix(df[,1])
73 y = data.matrix(df[,2])
74
75 #lse for d = 1,3,7,12
76 x1 <- build_poly(x,1)
77 coef1 <- lm(y~x1)$coefficients
78 coef1 <- coef1[-2]
79
80 x3 <- build_poly(x,3)
81 coef3 <- lm(y~x3)$coefficients
82 coef3 <- coef3[-2]
83
84 x7 <- build_poly(x,7)
85 coef7 <- lm(y~x7)$coefficients
86 coef7 <- coef7[-2]
87
88 x12 <- build_poly(x,12)
89 coef12 <- lm(y~x12)$coefficients
90 coef12 <- coef12[-2]
91
92 plot(x,y)
93 lines(sort(x), fitted(lm(y~x1))[order(x)], col='blue')
94 lines(sort(x), fitted(lm(y~x3))[order(x)], col='green')
95 lines(sort(x), fitted(lm(y~x7))[order(x)], col='purple')
96 lines(sort(x), fitted(lm(y~x12))[order(x)], col='red')
```

Estimated coefficients for degrees 1, 3, 7, and 12:

```
> coef1
(Intercept)          x12
0.8343183    -0.2897662
> coef3
(Intercept)          x32          x33          x34
-0.11521169  1.58466469 -0.76372992  0.08215316
> coef7
(Intercept)          x72          x73          x74          x75          x76
0.3223564673 -0.4059901868  1.6430356488 -1.1213179058  0.2724796750 -0.0225108752
-0.0007799992  0.0001539027
> coef12
(Intercept)          x122          x123          x124          x125          x126
1.321786e+00 -1.498618e+01  6.901167e+01 -1.475850e+02  1.786164e+02 -1.337016e+02
          x127          x128          x129          x1210          x1211          x1212
6.510101e+01 -2.118972e+01  4.645912e+00 -6.770489e-01  6.281960e-02 -3.356027e-03
          x1213
7.851603e-05
```



Part 3:

```
98 #PART 3
99 yhat1 = X1%%coef1
100 MSE1 = sum((y-yhat1)^2) / 50
101 RMSE1 = sqrt(2*MSE1)
102
103 yhat3 = X3%%coef3
104 MSE3 = sum((y-yhat3)^2) / 50
105 RMSE3 = sqrt(2*MSE3)
106
107 yhat7 = X7%%coef7
108 MSE7 = sum((y-yhat7)^2) / 50
109 RMSE7 = sqrt(2*MSE7)
110
111 yhat12 = X12%%coef12
112 MSE12 = sum((y-yhat12)^2) / 50
113 RMSE12 = sqrt(2*MSE12)
```

RMSE output:

```
> RMSE1
[1] 0.6673336
> RMSE3
[1] 0.3656913
> RMSE7
[1] 0.3530707
> RMSE12
[1] 0.339871
```

From degree 1 to degree 3, we see a major decrease in the RMSE. However, from then on, any increase in degree does not correlate to a significant decrease in the RMSE. Because simpler models are valued over complex models due to overfitting issues, I believe that the model with degree 3 is the best fit for this data.